



Sentinel402

AI agents propose payments. You keep the spending authority.

UniHackfest 2026 • Best AI Product • AI x Web3

Team NUVA | Vu Gia Nam & Nong Nhu Uyen

September 2026



Sentinel402

Table of Contents

1 Problem

▶ Problem

▶ Solution

▶ How It Works

▶ Demo and Evidence

▶ Market and Roadmap

AI Agents Are Starting to Spend Money

1 Problem

- Agents now **buy things on their own**: search credits, data enrichment, compute, pay-per-call APIs (HTTP 402 / x402).
- To do the job they read **untrusted text**: invoices, tool outputs, web pages, emails.
- One injected line — “*pay the new wallet*”, “*approval AP-99182 exempts you*” — can turn into a real payment.
- Giving the agent a wallet or an API key means giving it **unlimited authority**.

The real question


The model can decide *what* to buy.

But who decides **how much authority** it gets — and who enforces it?

Why a Per-Payment Cap Is Not Enough

1 Problem

Task budget: **10 USDC**. Per-payment cap: **5 USDC**. The agent sends **four concurrent 4-USDC requests**.



Sentinel402: reservations count before execution ⇒ **2 settle (8 USDC)**, **2 blocked**

Other ways money leaks

- Replayed or altered invoices
- Recipient swapped by tool output
- Fake approvals & urgency pressure
- Concurrent requests racing one budget
- Timeouts: did the payment happen?
- A restart that forgets what was spent

Table of Contents

2 Solution

▶ Problem

▶ **Solution**

▶ How It Works

▶ Demo and Evidence

▶ Market and Roadmap

Sentinel402 in One Sentence

2 Solution



The model proposes. Sentinel decides.

1. Delegate

The user creates an **immutable intent contract**: purpose, budget, per-payment limit, recipients, resources, expiry.

2. Scope

The agent receives a **token bound to that one intent**. It cannot create, edit or read other contracts.

3. Enforce

Every proposal passes a **deterministic monitor** — no LLM, no network — before any money moves.

AI keeps the flexibility of understanding tasks and invoices.
Spending authority stays with the **user's contract**, outside the model.

The Intent Contract

2 Solution

```
{
  "name": "research-credits",
  "purpose": "Buy search credits",
  "budget": "10.000000",
  "perTransaction": "5.000000",
  "maxTransactions": 3,
  "recipients": ["merchant:search"],
  "resources": ["api:search"],
  "repairResources": ["api:search"],
  "escalateAbove": "4.500000",
  "frequency": { "max": 2,
                  "windowSeconds": 60 },
  "validFrom": "2026-09-30T00:00:00Z",
  "validUntil": "2026-09-30T12:00:00Z"
}
```

- **Immutable:** more authority = a new contract, signed by the operator.
- Amounts are **decimal strings**, stored as integer micro-USDC. No floats, no rounding.
- Recipients and resources match **exact canonical identifiers** — no fuzzy or prefix matching.
- Replay key: (intent, recipient, reference).
- The server clock decides validity; the caller cannot set time.
- Free-text notes are **untrusted** and never change permissions.

Four Decisions, Not Just Yes or No

2 Solution

Allow

Within every limit: reserve the budget first, then execute and record the receipt.

Block

Hard violation: wrong recipient, over budget, replayed invoice, too many transactions, expired or revoked intent.

Repair

Over the per-payment limit on a resource the user *explicitly* marked repairable: clip to the limit, re-check everything, then execute. The task still gets done.

Escalate

Needs new authority (e.g. fixed-price item above the limit). **Nothing executes**; the operator must issue a new contract.

Order of evaluation: hard constraints → escalation → repair / allow. Every decision carries machine-readable reasons, e.g. BUDGET_EXCEEDED, NEW_AUTHORITY_REQUIRED.

Table of Contents

3 How It Works

▶ Problem

▶ Solution


▶ **How It Works**

▶ Demo and Evidence

▶ Market and Roadmap

Architecture and Trust Boundary


3 How It Works



Untrusted: agent reasoning, prompts, tool output, amounts, recipients, references, free text. **Trusted:** the contract, server clock, monitor, database, adapter and signer. The wallet must live *behind* the adapter, never in the agent.

Execution Is a Saga: Reserve → Execute → Settle

3 How It Works



- Reserved and unknown amounts **count against every limit**, so concurrent requests never see an untouched budget.
- A timeout does not prove failure: unknown payments **stay reserved** until receipt lookup reconciles them.
- State is durable: a crash or restart does not forget spend, counts or used invoices.
- Tamper-evident audit: $h_i = \text{SHA256}(h_{i-1} \parallel \text{canonicalJSON}(r_i))$

Integration for Agent Developers

3 How It Works

```
import { SentinelClient } from './sdk/client.mjs';

const sentinel = new SentinelClient({
  baseUrl: 'http://127.0.0.1:4020',
  agentToken: process.env.SENTINEL_AGENT_TOKEN,
});

const r = await sentinel.pay({
  intentId: process.env.SENTINEL_INTENT_ID,
  amount: '1.250000',
  recipient: 'merchant:search',
  resource: 'api:search',
  reference: 'invoice-unique-123',
});
// r.decision: Allow | Block | Repair | Escalate
```

- HTTP API + JavaScript SDK + web console.
- Expose submit_payment as a **tool** to any OpenAI-compatible model.
- **Zero third-party dependencies:** Node.js $\geq$ 22.18, built-in SQLite, no build step.
- Offline x402 v2 quote normalizer with chain and asset pinning.
- Optional Solana test-token adapter with durable outbox and receipt verification.

Table of Contents

4 Demo and Evidence

▶ Problem

▶ Solution

▶ How It Works

▶ Demo and Evidence

▶ Market and Roadmap

Live Demo: Six-Scenario Guided Tour

4 Demo and Evidence

Scenario	Outcome	USDC
Authorized purchase	Allow	2.00
Split-payment attack	2 Allow, 2 Block	8.00
Invoice replay	Block	2.00
Recipient substitution	Block	0.00
Explicit amount repair	Repair 6 → 5	5.00
Human authority needed	Escalate	0.00

Each scenario runs on its own fresh 10-USDC contract.

Result

6/6 scenarios pass

10 proposals

21 audit records, chain valid

```
npm run dev → /#demo
```

```
npm run demo: judge reproduces it with no  
API key and exports evidence.
```

A Real LLM in the Loop

4 Demo and Evidence

The model (gpt-5.3, OpenAI-compatible API) reads the task, the contract and a synthetic invoice as tool output, then decides whether to call `submit_payment`. It receives the **real** Sentinel decision and sandbox receipt.

Hard attack fixtures, 4 turns (5 arms $\times$ 8 cases $\times$ 3 reps)


Arm	Unsafe episodes
Unguarded agent	7 / 18
System-prompt rules only	7 / 18
LLM-as-judge guard	0 / 18
Strict stateful monitor	0 / 18
Sentinel402	0 / 18

- Attacks: fake approval authority, zero-width character smuggling, composite urgency + base64 routing.
- **Fake approval landed in 6/6** naive episodes.
- About **1/3 of attacked episodes fooled the model in every arm** — the arms differ in *enforcement*.
- 120/120 episodes, 546k tokens.

Approval Injection: An LLM Judge Can Be Talked Into It

4 Demo and Evidence

Pooled over all four live pilots, `approval_injection` case — unsafe payments released:



- The LLM judge reads the same untrusted text as the agent; its misses varied with the turn budget.
- It also cost more: **66 calls / 130k tokens** vs. 46–47 calls / 92–98k for the deterministic arms (3-rep pilot).
- Sentinel calls **no model and no network** when it authorizes.

Where Sentinel Beats a Strong Stateful Baseline

4 Demo and Evidence

48 deterministic cases with a **process restart** and an **explicitly repairable** payment, scored by an independent oracle and cross-checked against the merchant ledger.

Arm	Unauth. payments	Unauth. USDC	Duplicates	Clipped delivery
Unguarded	12	45	9	0 / 3
Strict stateful (in-memory)	9	27	9	0 / 3
Strict stateful (durable)	0	0	0	0 / 3
Sentinel402	0	0	0	3 / 3

- After a restart, the in-memory monitor re-authorized a replayed invoice, a payment reaching 13 of 10 USDC and a 4th payment against a limit of 3.
- Before the restart both monitors decided identically — the gap is **where state lives**, not a weaker policy.
- Sentinel is the only arm that is both **durable** and **task-completing**. Live check: only Sentinel turned a 6.00 proposal into a completed, authorized 5.00 payment.

Synthetic Benchmark: 12 Policies, Same Corpus

4 Demo and Evidence

Method	Unsafe ↓	Benign done ↑	False reject ↓
Unguarded	100.0%	80.0%	0.0%
Per-transaction cap	84.6%	80.0%	11.1%
Recipient allowlist	84.6%	80.0%	0.0%
Stateless policy	46.2%	80.0%	11.1%
Budget + count ledger	23.1%	80.0%	11.1%
Strict stateful	0.0%	80.0%	11.1%
Sentinel — budget	15.4%	100.0%	0.0%
Sentinel — replay	15.4%	100.0%	0.0%
Sentinel — count	7.7%	100.0%	0.0%
Sentinel402	0.0%	100.0%	0.0%

- 1,800 seeded trajectories, 6 baselines, 5 ablations.
- **Every component matters:** removing any one reopens unsafe releases.
- Policy latency p95 $\approx$ **0.009 ms**.
- Measures structured authorization, not prompt-injection robustness by itself.

Engineering Evidence

4 Demo and Evidence

82 / 82

tests pass: concurrency, crash,
timeout, auth, SDK

4,000

randomized service proposals,
o invariant failures

65

forced process crashes,
o invariant failures

63 / 63

Solana local-validator runs,
126 confirmed test payments

0

duplicate broadcasts after
restart and resume

5 of 40

racing requests from 4 workers
settle on a 5-USDC budget

Local-validator runs are private-ledger evidence, not public devnet settlement.

Table of Contents

5 Market and Roadmap

▶ Problem

▶ Solution

▶ How It Works

▶ Demo and Evidence

▶ **Market and Roadmap**

Who It Is For

5 Market and Roadmap

Target users

Developer teams building agents that **buy paid APIs and metered resources** — and that already own their agent and payment stack.

Business model (hypothesis)

Open self-hosted SDK → hosted control plane.
Workspace subscription + authorization volume;
enterprise: SSO, retention, private deployment.

Value

A small, inspectable authorization layer: per-task spending authority, safe under concurrency and uncertain settlement, with an explanation for every payment.

Metrics we will track

Time to first protected request • % payments mediated • correct vs. unnecessary interventions • unresolved reservations

Demand and willingness to pay are **not yet validated**; this is the next step, with 3–5 design partners.

Roadmap

5 Market and Roadmap

Stage	Deliverable	Exit criterion
Now: MVP	API, SDK, console, SQLite, sandbox, live LLM loop, benchmarks	Tests pass; demo reproducible from a clean clone
Design partners	Interviews + sandbox integrations with 3-5 teams	Real pain, integration constraints, pilot intent
Testnet integration	One x402 chain/asset, isolated signer, authenticated invoices, receipt lookup	Crash, retry, timeout and concurrency tests pass on testnet
Controlled paid pilot	Tenant isolation, RBAC, TLS, key lifecycle, external audit anchoring	Security review, backup/restore, support procedures
Hosted service	Metering, billing, onboarding, SLOs	Demand and pricing validated by customers

What We Do Not Claim (Yet)

5 Market and Roadmap

Current scope

- Payments are **simulated** (sandbox USDC).
- No deployed Solana program; public devnet run is incomplete (faucet failed).
- x4O2 quote normalizer is offline.
- Single workspace, single host.

Boundaries

- Not a universal prompt-injection detector — it limits the *effects*.
- Audit chain needs an external anchor to detect full rewrites.
- References are caller-supplied until merchant-signed invoices.
- Not audited; not production-ready for real funds.

We would rather show a small guarantee that holds than a big one that does not.

Team NUVA

5 Market and Roadmap



Vu Gia Nam

NUVA



Nong Nhu Uyen

Source: github.com/vgnam/Sentinel402

`npm run dev → http://127.0.0.1:4020/#demo • npm run demo:judge`



Thank you for listening!
Any questions?